


“Only a brain-damaged operating system would support task switching and not make the simple next step of supporting multitasking.”

– Calvin Keegan


Processes

- Abstraction of a running program
- Unit of work in the system
- Pseudoparallelism
- A process is *traced* by listing the sequence of instructions that execute for that process
- The process model
 - Sequential Process/Task
 - * A program in execution
 - * Program code
 - * Current activity
 - * Process stack
 - subroutine parameters
 - return addresses
 - temporary variables
 - * Data section
 - Global variables
- Concurrent Processes
 - Multiprogramming
 - Interleaving of traces of different processes characterizes the behavior of the CPU
 - Physical resource sharing
 - * Required due to limited hardware resources
 - Logical resource sharing
 - * Concurrent access to the same resource like files
 - Computation speedup
 - * Break each task into subtasks
 - * Execute each subtask on separate processing element
 - Modularity
 - * Division of system functions into separate modules
 - Convenience
 - * Perform a number of tasks in parallel
 - Real-time requirements for I/O
- Process Hierarchies
 - Parent-child relationship
 - `fork(2)` call in Unix
 - In MS-DOS, parent suspends itself and lets the child execute
- Process states


- Running
- Ready (Not running, waiting for the CPU)
- Blocked / Wait on an event (other than CPU) (Not running)
- Two other states complete the five-state model – New and Exit
 - * A process being created can be said to be in state New; it will be in state Ready after it has been created
 - * A process being terminated can be said to be in state Exit



- Above model suffices for most of the discussion on process management in operating systems; however, it is limited in the sense that the system screeches to a halt (even in the model) if all the processes are resident in memory and they all are waiting for some event to happen
- Create a new state Suspend to keep track of blocked processes that have been temporarily kicked out of memory to make room for new processes to come in
- The state transition diagram in the revised model is



- Which process to grant the CPU when the current process is swapped out?
 - * Preference for a previously suspended process over a new process to avoid increasing the total load on the system
 - * Suspended processes are actually blocked at the time of suspension and making them ready will just change their state back to blocked
 - * Decide whether the process is blocked on an event (suspended or not) or whether the process has been swapped out (suspended or not)
- The new state transition diagram is



Process control

- Modes of execution
 - OS execution vs user process execution
 - OS may prevent execution of some instructions in user mode and allow them to be executed only in privileged mode (also called kernel mode, system mode, or control mode)
 - * Read/write a control register, such as PSW
 - * Primitive I/O and memory management
 - The two modes protect the OS data structures from interference by user code
 - Kernel mode provides full control of the system that may not be needed for user programs
 - The kernel mode can be entered by setting a bit in the PSW
 - The system can enter privileged mode as a result of a request from user code and returns to user mode after completing the request
- Implementation of processes
 - Process table
 - * One entry for each process
 - * program counter
 - * stack pointer
 - * memory allocation
 - * open files
 - * accounting and scheduling information
 - *Interrupt vector*
 - * Contains address of *interrupt service procedure*
 - saves all registers in the process table entry
 - services the interrupt
- Process creation
 - Assign a unique process identifier to the new process; add this process to the system process table that contains one entry for each process
 - Allocate space for all elements of process image – space for code, data, and user stack; values can be set by default or based on parameters entered at job creation time
 - Allocation of resources (CPU time, memory, files) – use either of the following policies
 - * New process obtains resources directly from the OS
 - * New process constrained to share resources from a subset of the parent process
 - Build the data structures that are needed to manage the process, especially process control block
 - When is a process created? – job submission, login, application such as printing
 - Static or dynamic process creation
 - Initialization data (input)
 - Process execution
 - * Parent continues to execute concurrently with its children
 - * Parent waits until all its children have terminated
- Process switching
 - Interrupt a running process and assign control to a different process
 - Difference between process switching and mode switching

- When to switch processes
 - * Any time when the OS has control of the system
 - * OS can acquire control by
 - Interrupt – asynchronous external event; not dependent on instructions; clock interrupt
 - Trap – Exception handling; associated with current instruction execution
 - Supervisor call – Explicit call to OS
- Processes in Unix
 - Identified by a unique integer – *process identifier*
 - Created by the **fork(2)** system call
 - * Copy the three segments (instructions, user-data, and system-data) without initialization from a program
 - * New process is the copy of the address space of the original process to allow easy communication of the parent process with its child
 - * Both processes continue execution at the instruction after the **fork**
 - * Return code for the **fork** is
 - zero for the child process
 - process id of the child for the parent process
 - Use **exec(2)** system call after **fork** to replace the child process's memory space with a new program (binary file)
 - * Overlay the image of a program onto the running process
 - * Reinitialize a process from a designated program
 - * Program changes while the process remains
 - **exit(2)** system call
 - * Finish executing a process
 - **wait(2)** system call
 - * Wait for child process to stop or terminate
 - * Synchronize process execution with the **exit** of a previously **forked** process
 - **brk(2)** system call
 - * Change the amount of space allocated for the calling process's data segment
 - * Control the size of memory allocated to a process
 - **signal(3)** library function
 - * Control process response to extraordinary events
 - * The complete family of **signal** functions (see man page) provides for simplified signal management for application processes
 - Daemons
 - * Background processes to do useful work on behalf of the user
 - Just sit in the machine, doing one or the other thing
 - * Differ from normal processes in the sense that daemons do not have a **stdin** or **stdout**, and sleep most of the time
 - Communication with humans achieved via logs
 - * Common daemons are
 - **update** to synchronize the file system with its image in kernel memory
 - **cron** for general purpose task scheduling
 - **lpd** or **lpsched** as a line printer daemon to pick up files scheduled for printing and distributing them to the printers

- `init` – the boss of it all
- `swapper` to handle kernel requests to swap pages of memory to/from disk

- MS-DOS Processes

- Created by a system call to load a specified binary file into memory and execute it
- Parent is suspended and waits for child to finish execution

- Process Termination

- Normal termination
 - * Process terminates when it executes its last statement
 - * Upon termination, the OS deletes the process
 - * Process may return data (output) to its parent
- Termination by another process
 - * Termination by the system call `abort`
 - * Usually terminated only by the parent of the process because
 - child may exceed the usage of its allocated resources
 - task assigned to the child is no longer required
- Cascading termination
 - * Upon termination of parent process
 - * Initiated by the OS

- `cobegin/coend`

- Also known as `parbegin/parend`
- Explicitly specify a set of program segments to be executed concurrently

```
cobegin
  p_1;
  p_2;
  ...
  p_n;
coend;
```

$$(a + b) \times (c + d) - (e/f)$$

```
cobegin
  t_1 = a + b;
  t_2 = c + d;
  t_3 = e / f;
coend
t_4 = t_1 * t_2;
t_5 = t_4 - t_3;
```

- `fork`, `join`, and `quit` Primitives

- More general than `cobegin/coend`
- `fork x`
 - * Creates a new process `q` when executed by process `p`
 - * Starts execution of process `q` at instruction labeled `x`
 - * Process `p` executes at the instruction following the `fork`
- `quit`

- * Terminates the process that executes this command
- join `t, y`
 - * Provides an indivisible instruction
 - * Provides the equivalent of test-and-set instruction in a concurrent language


```
if ( ! --t ) goto y;
```
- Program segment with new primitives


```

m = 3;
fork p2;
fork p3;
p1 : t1 = a + b; join m, p4; quit;
p2 : t2 = c + d; join m, p4; quit;
p3 : t3 = e / f; join m, p4; quit;
p4 : t4 = t1 × t2;
      t5 = t4 - t3;
```

Process Control Subsystem in Unix

- Significant part of the Unix kernel (along with the file subsystem)
- Contains three modules
 - Interprocess communication
 - Scheduler
 - Memory management

Interprocess Communication

- Race conditions
 - A race condition occurs when two processes (or threads) access the same variable/resource without doing any synchronization
 - One process is doing a coordinated update of several variables
 - The second process observing one or more of those variables will see inconsistent results
 - Final outcome dependent on the precise timing of two processes
 - Example
 - * One process is changing the balance in a bank account while another is simultaneously observing the account balance and the last activity date
 - * Now, consider the scenario where the process changing the balance gets interrupted after updating the last activity date but before updating the balance
 - * If the other process reads the data at this point, it does not get accurate information (either in the current or past time)

Critical Section Problem

- Section of code that modifies some memory/file/table while assuming its exclusive control
- Mutually exclusive execution in time
- Template for each process that involves critical section

```

do
{
    ...                /* Entry section;                */
    critical_section(); /* Assumed to be present          */
    ...                /* Exit section                    */
    remainder_section(); /* Assumed to be present          */
}
while ( 1 );

```

You are to fill in the gaps specified by ... for entry and exit sections in this template and test the resulting program for compliance with the protocol specified next

- Design of a protocol to be used by the processes to cooperate with following constraints
 - Mutual Exclusion – If process p_i is executing in its critical section, then no other processes can be executing in their critical sections.
 - Progress – If no process is executing in its critical section, the selection of a process that will be allowed to enter its critical section cannot be postponed indefinitely.
 - Bounded Waiting – There must exist a bound on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.
- Assumptions
 - No assumption about the hardware instructions
 - No assumption about the number of processors supported
 - Basic machine language instructions executed atomically
- Disabling interrupts
 - Brute-force approach
 - Not proper to give users the power to disable interrupts
 - * User may not enable interrupts after being done
 - * Multiple CPU configuration
- Lock variables
 - Share a variable that is set when a process is in its critical section
- Strict alternation

```

extern int turn;    /* Shared variable between both processes */

do
{
    while ( turn != i ) /* do nothing */ ;
    critical_section();
    turn = j;
    remainder_section();
} while ( 1 );

```

- Does not satisfy progress requirement
- Does not keep sufficient information about the state of each process
- Use of a flag

```

extern int flag[2];          /* Shared variable; one for each process */

do
{
    flag[i] = 1;              /* true */
    while ( flag[j] );
    critical_section();
    flag[i] = 0;              /* false */
    remainder_section();
} while ( 1 );

```

- Satisfies the mutual exclusion requirement
- Does not satisfy the progress requirement

Time T_0 p_0 sets flag[0] to true
 Time T_1 p_1 sets flag[1] to true

Processes p_0 and p_1 loop forever in their respective while statements

- Critically dependent on the exact timing of two processes
- Switch the order of instructions in entry section
 - * No mutual exclusion

- Peterson's solution

- Combines the key ideas from the two earlier solutions

/* Code for process 0; similar code exists for process 1 */

```

extern int flag[2];          /* Shared variables */
extern int turn;             /* Shared variable */

void process_0()
{
    do
    {
        /* Entry section */
        flag[0] = true;      /* Raise my flag */
        turn = 1;            /* Cede turn to other process */
        while ( flag[1] && turn == 1 );

        critical_section();

        /* Exit section */
        flag[0] = false;

        remainder_section();

    } while ( 1 );
}

```

- Multiple Process Solution – Solution 4

- The array flag can take one of the three values (idle, want-in, in-cs)

```

enum state { idle, want_in, in_cs };
extern int turn;
extern state flag[n];    // Flag corresponding to each process (in shared memory)

```



```

// Code for process i

int  j;                // Local to each process

do
    do
        flag[i] = want_in;    // Raise my flag
        j = turn;            // Set local variable
        while ( j != i )
            j = ( flag[j] != idle ) ? turn : ( j + 1 ) % n;

        // Declare intention to enter critical section

        flag[i] = in_cs;

        // Check that no one else is in critical section

        for ( j = 0; j < n; j++ )
            if ( ( j != i ) && ( flag[j] == in_cs ) )
                break;

        while ( j < n ) || ( turn != i && flag[turn] != idle );

        // Assign turn to self and enter critical section

        turn = i;
        critical_section();

        // Exit section

        j = (turn + 1) % n;
        while (flag[j] == idle) do
            j = (j + 1) % n;

        // Assign turn to the next waiting process and change own flag to idle

        turn = j;
        flag[i] = idle;

        remainder_section();
    while ( 1 );

```

- p_i enters the critical section only if $\text{flag}[j] \neq \text{in_cs}$ for all $j \neq i$.
- **turn** can be modified only upon entry to and exit from the critical section. The first contending process enters its critical section.
- Upon exit, the successor process is designated to be the one following the current process.
- Mutual Exclusion
 - * p_i enters the critical section only if $\text{flag}[j] \neq \text{in_cs}$ for all $j \neq i$.
 - * Only p_i can set $\text{flag}[i] = \text{in_cs}$.
 - * p_i inspects $\text{flag}[j]$ only while $\text{flag}[i] = \text{in_cs}$.
- Progress
 - * **turn** can be modified only upon entry to and exit from the critical section.

- * No process is executing or leaving its critical section \Rightarrow `turn` remains constant.
- * First contending process in the cyclic ordering (`turn`, `turn+1`, ..., `n-1`, 0, ..., `turn-1`) enters its critical section.
- Bounded Wait
 - * Upon exit from the critical section, a process must designate its unique successor the first contending process in the cyclic ordering `turn+1`, ..., `n-1`, 0, ..., `turn-1`, `turn`.
 - * Any process waiting to enter its critical section will do so in at most `n-1` turns.

- Bakery Algorithm

- Each process has a unique id
- Process id is assigned in a completely ordered manner

```
extern bool choosing[n];    /* Shared Boolean array */
extern int  number[n];      /* Shared integer array to hold turn number */

void process_i ( const int i )    /* ith Process */
{
    do
    {
        choosing[i] = true;
        number[i] = 1 + max(number[0], ..., number[n-1]);
        choosing[i] = false;
        for ( int j = 0; j < n; j++ )
        {
            while ( choosing[j] );    /* Wait while someone else is choosing */
            while ( ( number[j] ) && (number[j],j) < (number[i],i) );
        }

        critical_section();

        number[i] = 0;

        remainder_section();
    } while ( 1 );
}
```

- If p_i is in its critical section and p_k ($k \neq i$) has already chosen its `number[k] \neq 0`, then `(number[i],i) < (number[k],k)`.

Synchronization Hardware

- `test_and_set` instruction

```
int test_and_set (int& target )
{
    int tmp;
    tmp = target;
    target = 1; /* True */
    return ( tmp );
}
```

- Implementing Mutual Exclusion with `test_and_set`

```

extern bool lock ( false );

do
    while ( test_and_set ( lock ) );
    critical_section();
    lock = false;
    remainder_section();
while ( 1 );

```

Semaphores

- Producer-consumer Problem
 - Shared buffer between producer and consumer
 - Number of items kept in the variable `count`
 - Printer spooler
 - The `|` operator
 - Race conditions
- An integer variable that can only be accessed through two standard atomic operations – wait (P) and signal (V)

Operation	Semaphore	Dutch	Meaning
Wait	P	<i>proberen</i>	test
Signal	V	<i>verhogen</i>	increment

- The classical definitions for *wait* and *signal* are

```

wait ( S ):    while ( S <= 0 );
               S--;

signal ( S ):  S++;

```

- Mutual exclusion implementation with semaphores

```

do
    wait (mutex);
    critical_section();
    signal (mutex);
    remainder_section();
while ( 1 );

```

- Synchronization of processes with semaphores

p_1	S_1 ; <code>signal (synch);</code>
p_2	<code>wait (synch);</code> S_2 ;

- Implementing Semaphore Operations
 - Binary semaphores using `test_and_set`
 - * Check out the instruction definition as previously given
 - Implementation with a *busy-wait*

```

class bin_semaphore
{
    private:
        bool        s;        /* Binary semaphore    */

    public:
        bin_semaphore()        // Default constructor
        : s ( false )
        {}

        void P()                // Wait on semaphore
        {
            while ( test_and_set ( s ) );
        }

        void V ()               // Signal the semaphore
        {
            s = false;
        }
};

```

– General semaphore

```

class semaphore
{
    private:
        bin_semaphore    mutex;
        bin_semaphore    delay;
        int               count;

    public:
        void semaphore ( const int num = 1 )    // Constructor
        : count ( num )
        {
            delay.P();
        }

        void P()
        {
            mutex.P();
            if ( --count < 0 )
            {
                mutex.V();
                delay.P();
            }
            mutex.V();
        }

        void V()
        {
            mutex.P();
            if ( ++count <= 0 )
                delay.V();
            else
                mutex.V();
        }
};

```

- ```

 }
}

```
- Busy-wait Problem – Processes waste CPU cycles while waiting to enter their critical sections
    - \* Modify **wait** operation into the **block** operation. The process can block itself rather than busy-waiting.
    - \* Place the process into a wait queue associated with the critical section
    - \* Modify **signal** operation into the **wakeup** operation.
    - \* Change the state of the process from *wait* to *ready*.
  - Block-Wakeup Protocol
- // Semaphore with block wakeup protocol

```

class sem_int
{
private:
 int value; // Number of resources
 queue<pid_t> l; // List of processes

public:
 void sem_int (const int n = 1) // Constructor
 : value (n)
 {
 l = queue<pid_t>(0); // Empty queue
 }

 void P()
 {
 if (--value < 0)
 {
 pid_t p = getpid();
 l.enqueue (p); // Enqueue the invoking process
 block (p);
 }
 }

 void V()
 {
 if (++value <= 0)
 {
 process p = l.dequeue();
 wakeup (p);
 }
 }
};

```

### Producer-Consumer problem with semaphores

```

extern semaphore mutex; // To get exclusive access to buffers
extern semaphore empty (n); // Number of available buffers
extern semaphore full (0); // Initialized to 0

void producer()
{
 do

```

```

{
 produce (item);
 empty.P(); // empty is semaphore
 mutex.P(); // mutex is semaphore
 put (item);
 mutex.V();
 full.V()
} while (1);
}

void consumer()
{
 do
 {
 full.P();
 mutex.P();
 remove (item);
 mutex.V();
 empty.V();
 consume (item);
 } while (1);
}

```

Problem: What if order of wait is reversed in producer

### Event Counters

- Solve the producer-consumer problem without requiring mutual exclusion
- Special kind of variable with three operations
  1. `E.read()`: Return the current value of E
  2. `E.advance()`: Atomically increment E by 1
  3. `E.await(v)`: Wait until E has a value of v or more
- Event counters always start at 0 and always increase

```

class event_counter
{
 int ec; // Event counter

public:
 event_counter () // Default constructor
 : ec (0)
 {}
 int read() const { return (ec); }
 void advance() { ec++; }
 void await (const int v) const { while (ec < v); }
};

extern event_counter in, out; // Shared event counters

void producer()
{
 int sequence (0); // Local to producer

```

```

do
{
 produce (item);
 sequence++;
 out.await (sequence - num_buffers);
 put (item);
 in.advance();
}
while (1);
}

void consumer()
{
 int sequence (0); // Local to consumer
 do
 {
 sequence++;
 in.await (sequence);
 remove (item);
 out.advance();
 consume (item);
 }
 while (1);
}

```

### Higher-Level Synchronization Methods

- P and V operations do not permit a segment of code to be designated explicitly as a critical section.
- Two parts of a semaphore operation; should be treated as distinct
  - Block-wakeup of processes
  - Counting of semaphore
- Possibility of a *deadlock* – Omission or unintentional execution of a V operation.
- Monitors
  - Implemented as a class with private and public functions
  - Collection of data [resources] and private functions to manipulate this data
  - A monitor must guarantee the following:
    - \* Access to the resource is possible only via one of the monitor procedures.
    - \* Procedures are mutually exclusive in time. Only one process at a time can be active within the monitor.
  - Additional mechanism for synchronization or communication – the **condition** construct
 

**condition x;**

    - \* **condition** variables are accessed by only two operations – **wait** and **signal**
    - \* **x.wait()** suspends the process that invokes this operation until another process invokes **x.signal()**
    - \* **x.signal()** resumes exactly one suspended process; it has no effect if no process is suspended
  - Selection of a process to execute within monitor after **signal**
    - \* **x.signal()** executed by process P allowing the suspended process Q to resume execution

1. P waits until Q leaves the monitor, or waits for another condition
  2. Q waits until P leaves the monitor, or waits for another condition
- Choice 1 advocated by Hoare

- The Dining Philosophers Problem – Solution by Monitors

```
enum state_type { thinking, hungry, eating };

class dining_philosophers
{
private:
 state_type state[5]; // State of five philosophers
 condition self[5]; // Condition object for synchronization

 void test (int i)
 {
 if ((state[(i + 4) % 5] != eating) &&
 (state[i] == hungry) &&
 (state[(i + 1) % 5] != eating))
 {
 state[i] = eating;
 self[i].signal();
 }
 }

public:
 void dining_philosophers() // Constructor
 {
 for (int i = 0; i < 5; state[i++] = thinking);
 }

 void pickup (const int i) // i corresponds to the philosopher
 {
 state[i] = hungry;
 test (i);
 if (state[i] != eating)
 self[i].wait();
 }

 void putdown (const int i) // i corresponds to the philosopher
 {
 state[i] = thinking;
 test ((i + 4) % 5);
 test ((i + 1) % 5);
 }
}
```

- Philosopher *i* must invoke the operations `pickup` and `putdown` on an instance `dp` of the `dining_philosophers` monitor

```
dining_philosophers dp;

dp.pickup(i); // Philosopher i picks up the chopsticks
...
dp.eat(i); // Philosopher i eats (for random amount of time)
```



```

...
dp.putdown(i); // Philosopher i puts down the chopsticks

```

- No two neighbors eating simultaneously – no deadlocks
- Possible for a philosopher to starve to death

- Implementation of a Monitor

- Execution of procedures must be mutually exclusive
- A **wait** must block the current process on the corresponding **condition**
- If no process is running in the monitor and some process is waiting, it must be selected. If more than one waiting process, some criterion for selecting one must be deployed.
- Implementation using semaphores

- \* Semaphore **mutex** corresponding to the monitor initialized to 1
    - Before entry, execute **wait(mutex)**
    - Upon exit, execute **signal(mutex)**
  - \* Semaphore **next** to suspend the processes unable to enter the monitor initialized to 0
  - \* Integer variable **next\_count** to count the number of processes waiting to enter the monitor
- ```

mutex.wait();

```

```

...
void P() { ... }    // Body of P()

```

```

...
if ( next_count > 0 )
    next.signal();
else
    mutex.signal();

```

- * Semaphore **x_sem** for condition **x**, initialized to 0
- * Integer variable **x_count**

```

class condition
{
    int          num_waiting_procs;    // Processes waiting on this condition
    semaphore    sem;                 // To synchronize the processes
    static int    next_count;          // Processes waiting to enter monitor
    static semaphore next;
    static semaphore mutex;

public:
    condition()           // Default constructor
    : num_waiting_procs ( 0 ), sem ( 0 )
    {}

    void wait()
    {
        num_waiting_procs++;          // # of processes waiting on this condition
        if ( next_count > 0 )          // Someone waiting inside monitor?
            next.signal();             // Yes, wake him up
        else
            mutex.signal();            // No, free mutex so others can enter
        sem.wait();                   // Start waiting for condition
        num_waiting_procs--;           // Wait over, decrement variable
    }
}

```

```

void signal()
{
    if ( num_waiting_procs <= 0 )    // Nobody waiting?
        return;
    next_count++;                    // Number of ready processes inside monitor
    sem.signal();                     // Send the signal
    next.wait();                      // You wait; let signalled process run
    next_count--;                    // One less process in monitor
}
};

```

- Conditional Critical Regions (CCRs)

- Designed by Hoare and Brinch-Hansen to overcome the deficiencies of semaphores
- Explicitly designate a portion of code to be critical section
- Specify the variables (resource) to be protected by the critical section

`resource r :: v_1, v_2, ..., v_n`

- Specify the conditions under which the critical section may be entered to access the elements that form the resource

`region r when B do S`

* B is a condition to guard entry into critical section S

* At any time, only one process is permitted to enter the code segment associated with resource r

- The statement `region r when B do S` is implemented by

```

semaphore mutex ( 1 ), delay ( 0 );
int          delay_cnt ( 0 );

mutex.P();
del_cnt++;
while ( !B )
{
    mutex.V();
    delay.P();
    mutex.P();
}
del_cnt--;
S;           // Critical section code
for ( int i ( 0 ); i < del_cnt; i++ )
    delay.V();
mutex.V();

```

Message-Based Synchronization Schemes

- Communication between processes is achieved by:

- Shared memory (semaphores, CCRs, monitors)
- Message systems
 - * Desirable to prevent sharing, possibly for security reasons or no shared memory availability due to different physical hardware

- Communication by Passing Messages

- Processes communicate without any need for shared variables

- Two basic communication primitives

- * send message

```
* receive message
```

```
send(P, message)
```

Send a message to process P

```
receive(Q, message)
```

Receive a message from process Q

- Messages passed through a *communication link*

- Producer/Consumer Problem

```
void producer ( void )
```

{

```
while ( 1 )
```

{

```
produce ( data );
```

```
send ( consumer, data );
```

}

}

```
void consumer ( void )
```

{

```
while ( 1 )
```

{

```
receive ( producer, data );
```

```
consume ( data );
```

}

}

- Issues to be resolved in message communication

- Synchronous v/s Asynchronous Communication

* Upon **send**, does the sending process continue (asynchronous or nonblocking communication), or does it wait for the message to be accepted by the receiving process (synchronous or blocking communication)?

* What happens when a **receive** is issued and there is no message waiting (blocking or nonblocking)?

- *Implicit v/s Explicit Naming*

- * Does the sender specify exactly one receiver (explicit naming) or does it transmit the message to all the other processes (implicit naming)?

send (p, message) Send a message to process p

send (A, message) Send a message to mailbox A

* Does the receiver accept from a certain sender (explicit naming) or can it accept from any sender (implicit naming)?

receive (p, message)	Receive a message from process p
----------------------	-------------------------------------

receive (id, message)	Receive a message from any process; id is the process id
-----------------------	--

receive (A, message)	Receive a message from mailbox A
----------------------	-------------------------------------

Ports and Mailboxes

- Achieve synchronization of asynchronous process by embedding a busy-wait loop, with a non-blocking `receive` to simulate the effect of implicit naming

- Inefficient solution

- Indirect communication avoids the inefficiency of busy-wait

- Make the queues holding messages between senders and receivers visible to the processes, in the form of mailboxes

- Messages are sent to and received from mailboxes

- Most general communication facility between n senders and m receivers
- Unique identification for each mailbox
- A process may communicate with another process by a number of different mailboxes
- Two processes may communicate only if they have a shared mailbox
- Properties of a communication link
 - A link is established between a pair of processes only if they have a shared mailbox
 - A link may be associated with more than two processes
 - Between each pair of communicating processes, there may be a number of different links, each corresponding to one mailbox
 - A link may be either unidirectional or bidirectional
- Ports
 - In a distributed environment, the **receive** referring to same mailbox may reside on different machines
 - Port is a limited form of mailbox associated with only one receiver
 - All messages originating with different processes but addressed to the same port are sent to one central place associated with the receiver

Remote Procedure Calls

- High-level concept for process communication, allowing functions to be called without using send/receive primitives
 - send/receive work like semaphores, taking attention away from the task at hand
 - RPCs allow the called function to be perceived as a service request
- Transfers control to another process, possibly on a different computer, while suspending the calling process
- Called procedure resides in separate address space and no global variables are shared
- Return statement executed by called function returns control to the caller
- Communication strictly by parameters

```
send (RP_guard, parameters);
receive (RP_guard, results);
```

- The remote procedure guard is implemented by

```
void RP_guard ( void )
{
    do
        receive (caller, parameters);
        ...
        send (caller, results);
    while ( 1 );
}
```

- Static versus dynamic creation of remote procedures
- rendezvous mechanism in Ada